

# 全链路监控那些事

黄杰@饿了么

# 大纲

- 介绍
- 整体架构
- 计算框架
- 存储方案
- Demo
- Q&A

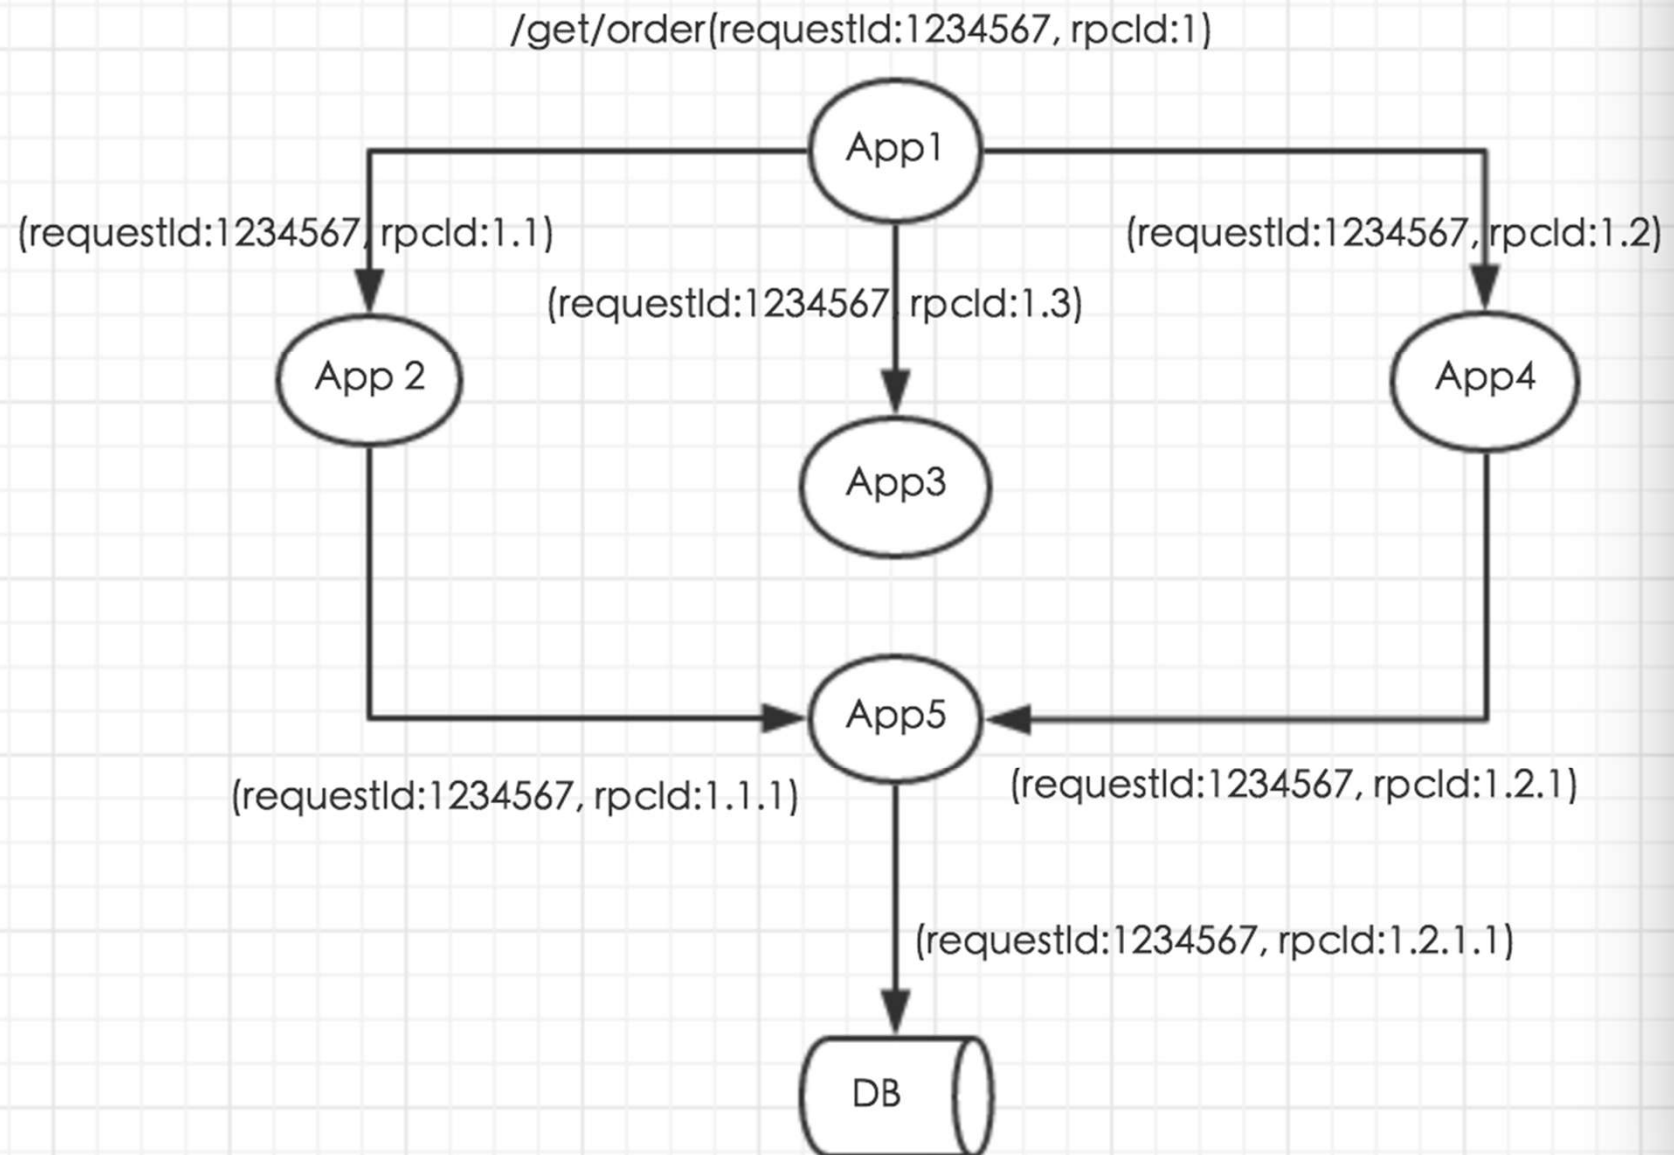
# 背景

- 从单应用到服务化之后，某个服务或几个服务出错时，不知道到底哪里出问题了；
- 应用大量的报错很多时间可能都被忽略；
- 有时应用出问题之后，需要花很长时间去定位问题；
- 服务之间的调用，数据库的访问等跨网络的IO操作，对我们都是黑盒子；
- 访问好慢，但不知道是哪里导致访问慢了；
- 对应具体某个应用内部也是黑盒的；

# 介绍

- 基于Java开发的轻量分布式实时监控平台，支持水平扩展，高吞吐；
- 由于饿了么开发语言的多样化，需要支持Java/Python/Go/Node.js；
- 支持跨IDC链路监控；
- 全量数据收集及处理；
- 由于是一个监控系统，所以整个系统不是一个高可靠的系统；

# 介绍

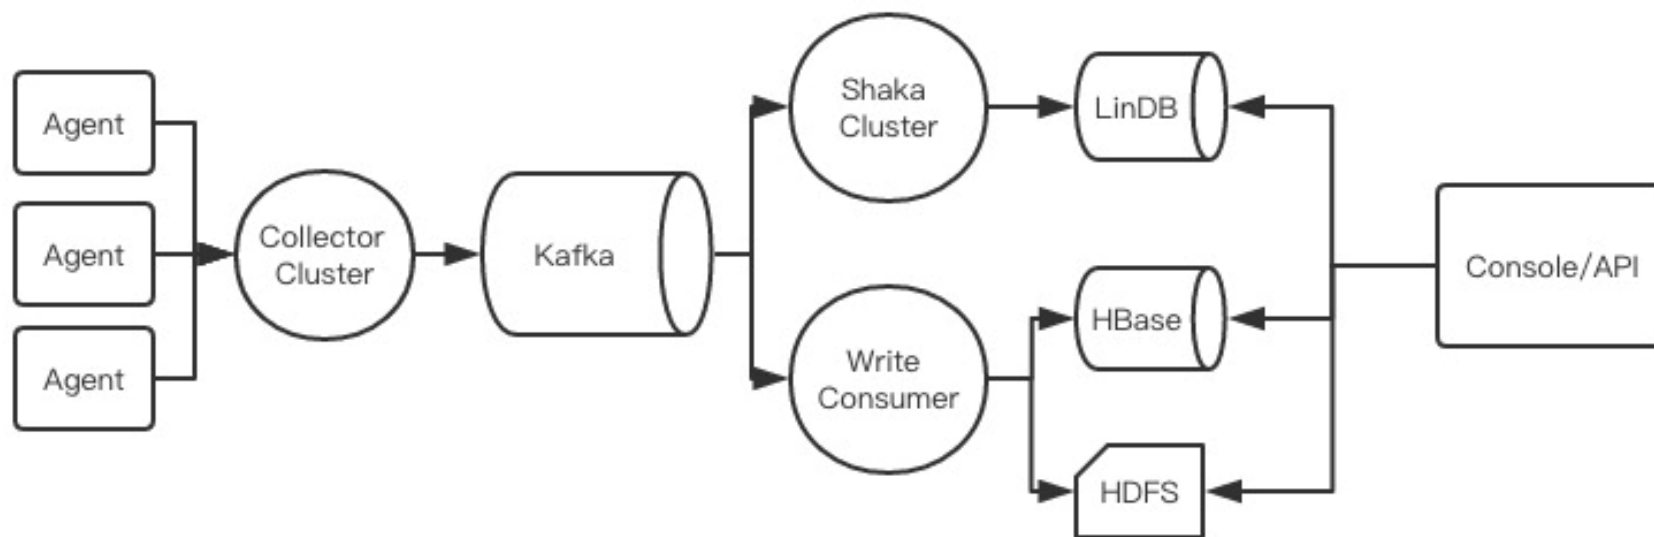




# 介绍

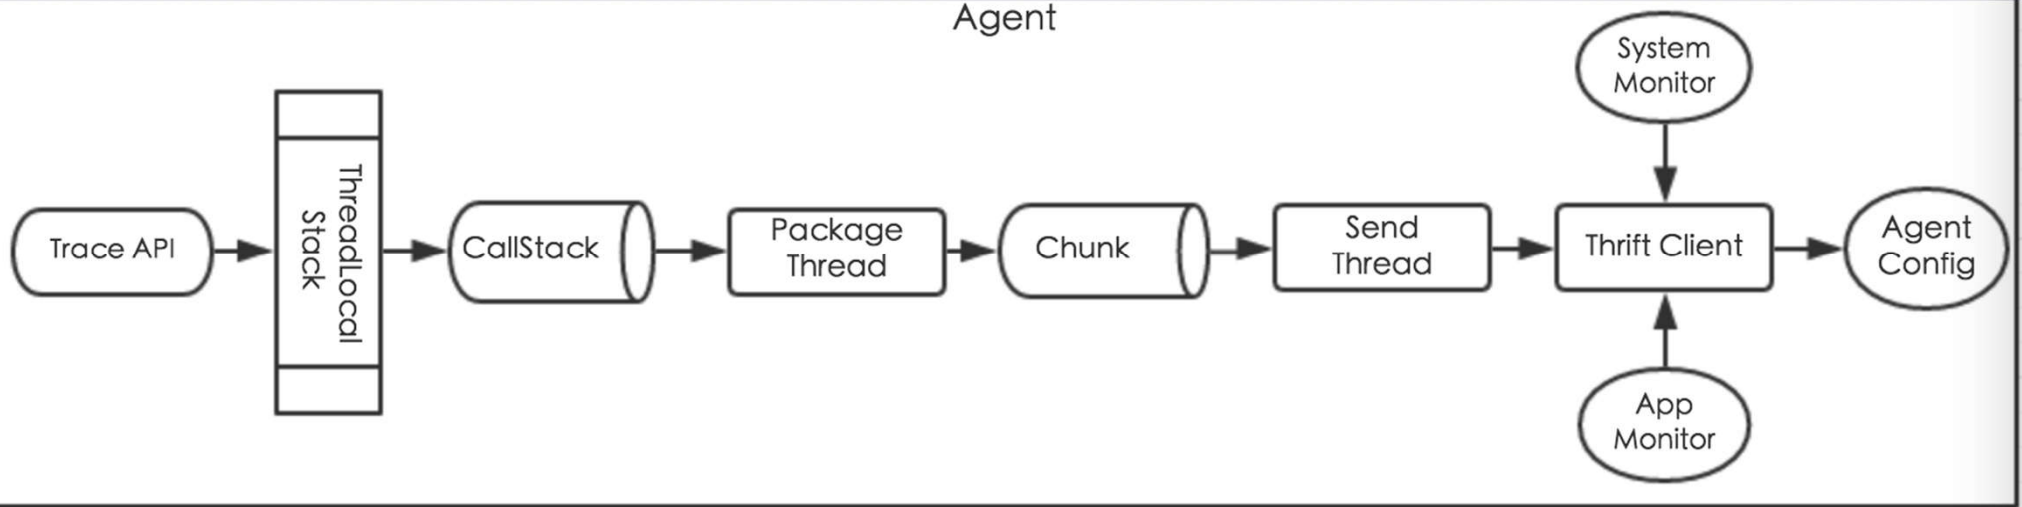
- 每天增量>200T, 压缩之后>60T;
- 接入应用>1000;
- 接入服务器>12000台, 不包括Docker;
- 消息处理延时<1s;
- 峰值消息处理QPS>2000W;

# 整体架构



# Agent

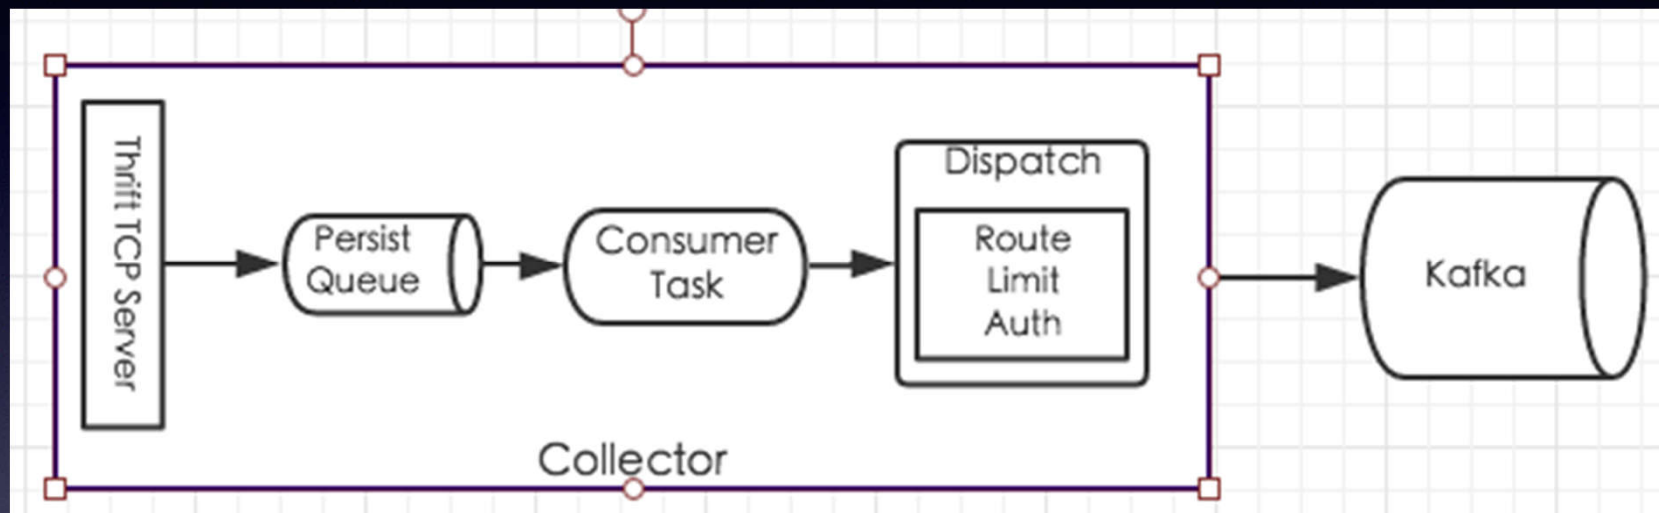
Agent



Duration	Duration(%)	Description(Type:Name)	ELog	EZone	Drill Up/Down	Timestamp
903.000ms	100.00%	SOAService:ShoppingService	Elog	wg1	Drill Up	2017-12-18 11:29:47.569
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:47.573
188.000ms	20.82%	SOACall:RankService.sea			Drill Down	2017-12-18 11:29:47.575
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:47.773
13.000ms	1.44%	SOACall:ShopBroadcastInfoSer			Drill Down	2017-12-18 11:29:47.774
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:47.808
1.000ms	0.11%	SOACall:ABTestService.getFactors			Drill Down	2017-12-18 11:29:47.810
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:47.826
49.000ms	5.43%	SOACall:ElemePromotionService			Drill Down	2017-12-18 11:29:47.828
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:47.878
338.000ms	37.21%	SOACall:ElemePromotionSe			Drill Down	2017-12-18 11:29:47.879
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:48.215
2.000ms	0.22%	SOACall:MemberSe			Drill Down	2017-12-18 11:29:48.217
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:48.219
13.000ms	1.44%	SOACall:ElemePromotionSe			Drill Down	2017-12-18 11:29:48.220
11.000ms	1.22%	Client:Getclient				2017-12-18 11:29:48.233
158.000ms	17.50%	SOACall:ElemePromotion			Drill Down	2017-12-18 11:29:48.245
1.000ms	0.11%	Client:Getclient				2017-12-18 11:29:48.404
49.000ms	5.43%	SOACall:ElemePromotionSen			Drill Down	2017-12-18 11:29:48.406
9.000ms	1.00%	Redis:Stats				2017-12-18 11:29:48.472



# Collector



# 计算框架

- 目前的数据处理及计算是基本Esper自己实现的;
- 没有使用目前比较流行的Storm/Spark Streaming这个的架构;
- 本身的业务逻辑不是很复杂, 引入Storm这样的组件同时也增加了运维成本, Storm之身的一些问题, 如雪崩等;
- 网络资源比较浪费;

# Shaka

Policy Config ×

Channel name:

app\_agg\_altb

\*Name:

exception

Desc:

Predefine event:

≡ event

≡ exception

≡ jvm

≡ redis

≡ soa\_call

≡ soa\_service

≡ transaction

≡ url

Time Window:

\*EPL expression:

```
1 @Metrics(prefix="appId",tags={"name","type","ezone","hostName"})
2 context context_name_1minute
3 select appId,name,type,ezone,hostName,timeMinutes as timestamp,counter(id) as count from exception
4 group by appId,name,type,ezone,timeMinutes,hostName
5 output snapshot when terminated
```

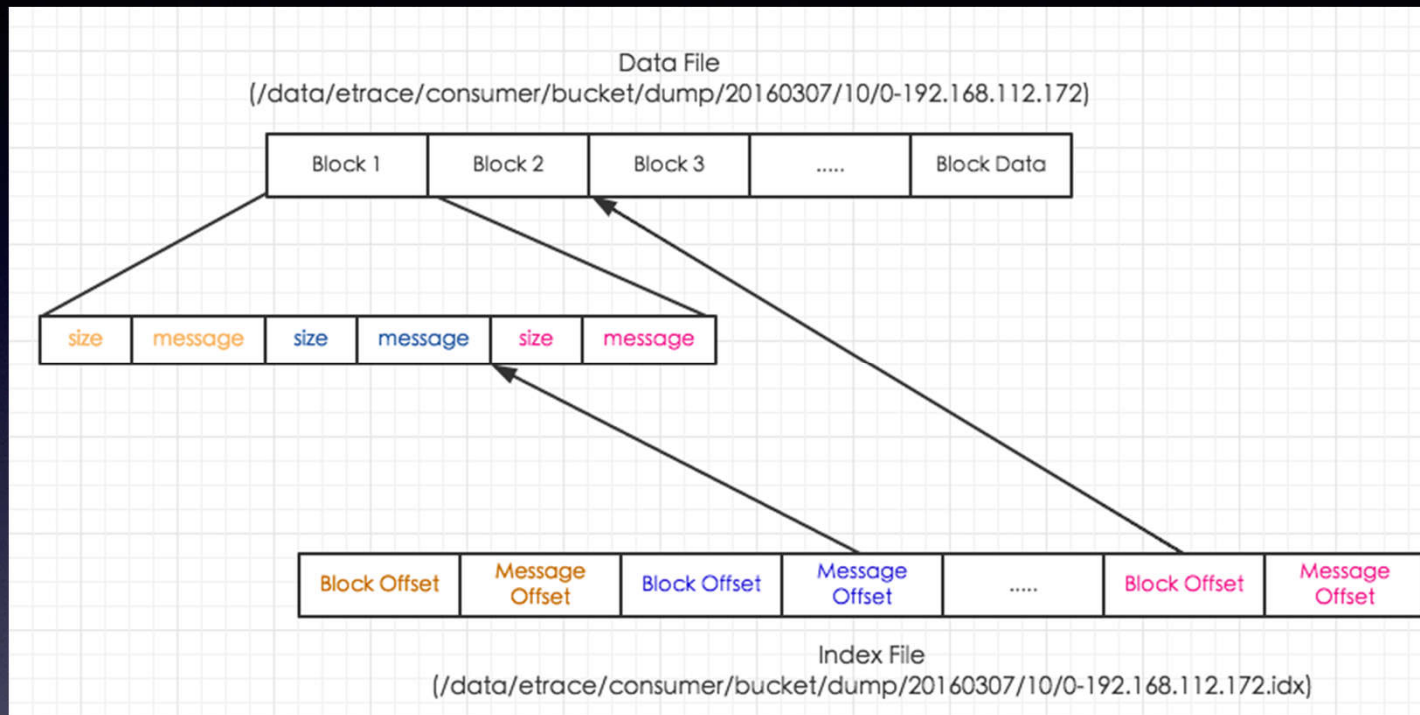
? Functions

Save

# 存储方案

- 存储使用HDFS及HBase，而不是目前比较火的ELK；
- HDFS存储Raw Data；
- HBase存储Index数据及一些采样数据；
- Raw Data按Block存储，每个Block通过Snappy压缩；
- 整个存储过程使用多个异步进行处理；
- 所有的指标目前存储在LinDB中；

# File Format



- 在Block写入的时候，为每个Block生成一个递增的Id；
- 每个Block 64K + Snappy Compress；
- Index File存储Block Offset(long)+Message Offset(long)；



# HBase

- 目前在HBase中存储index和采样数据;
- HBase主要解决热点问题, 及Key的设计是否合理;
- 通过pre-split region, 把region分散在每个Region Server中;
- 写入的时候根据Region把数据分散在每个Region上面;

# Demo

Overview

RuntimeException 3.03K 11.9 31%

BusinessException 63.37K 1.18 22%

Exception 165.81K

机器分布图

Trace Detail

Topology

Duration	Duration(%)	Description(Type:Name)	ELog	EZone	Drill Up/Down	Timestamp
7.000ms	100.00%	SOA\$	Elog	xg1	Drill Up	2017-12-18 12:44:39.604
6.000ms	85.71%	PaymentS				2017-12-18 12:44:39.605
1.000ms	14.29%	SQL			Drill Down	2017-12-18 12:44:39.606

Status: com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationExceptionTimestamp: 2017-12-18 12:44:39.606

2017-12-18 12:44:39.39] - com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationException  
com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationException: Duplicate entry '5-1215045249770997798-0' for key 'out\_trade\_no\_unique\_constraint'  
at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)  
at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)  
at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)  
at java.lang.reflect.Constructor.newInstance(Constructor.java:423)  
at com.mysql.jdbc.Util.handleNewInstance(Util.java:389)  
at com.mysql.jdbc.Util.getInstance(Util.java:372)  
at com.mysql.jdbc.SQLError.createSQLException(SQLError.java:973)  
at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3835)  
at com.mysql.jdbc.MysqlIO.checkErrorPacket(MysqlIO.java:3771)  
at com.mysql.jdbc.MysqlIO.sendCommand(MysqlIO.java:2435)  
at com.mysql.jdbc.MysqlIO.sqlQueryDirect(MysqlIO.java:2582)  
at com.mysql.jdbc.ConnectionImpl.execSQL(ConnectionImpl.java:2535)  
at com.mysql.jdbc.PreparedStatement.executeInternal(PreparedStatement.java:1911)  
at com.mysql.jdbc.PreparedStatement.execute(PreparedStatement.java:1203)  
at me.ele.ejdbc.impl.EPreparedStatement.loggedExecute(EPreparedStatement.java:79)  
at me.ele.ejdbc.impl.EPreparedStatement.execute(EPreparedStatement.java:111)  
at com.mchange.v2.c3p0.impl.NewProxyPreparedStatement.execute(NewProxyPreparedStatement.java:989)  
at org.apache.ibatis.executor.statement.PreparedStatementHandler.update(PreparedStatementHandler.java:44)  
at org.apache.ibatis.executor.statement.RoutingStatementHandler.update(RoutingStatementHandler.java:69)  
at org.apache.ibatis.executor.SimpleExecutor.doUpdate(SimpleExecutor.java:48)  
at org.apache.ibatis.executor.BaseExecutor.update(BaseExecutor.java:105)  
at org.apache.ibatis.executor.CachingExecutor.update(CachingExecutor.java:71)  
at org.apache.ibatis.session.defaults.DefaultSqlSession.update(DefaultSqlSession.java:152)  
at org.apache.ibatis.session.defaults.DefaultSqlSession.insert(DefaultSqlSession.java:141)  
at sun.reflect.GeneratedMethodAccessor366.invoke(Unknown Source)  
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)  
at java.lang.reflect.Method.invoke(Method.java:498)  
at org.mybatis.spring.SqlSessionTemplate\$SqlSessionInterceptor.invoke(SqlSessionTemplate.java:358)  
at com.sun.proxy.\$Proxy23.insert(Unknown Source)  
at org.mybatis.spring.SqlSessionTemplate.insert(SqlSessionTemplate.java:240)  
at org.apache.ibatis.binding.MapperMethod.execute(MapperMethod.java:51)  
at org.apache.ibatis.binding.MapperProxy.invoke(MapperProxy.java:52)  
at com.sun.proxy.\$Proxy78.insertSelective(Unknown Source)  
at me.ele... java:139)  
at me.ele... ceImpl.java:120)  
at me.ele... java:189)  
at me.ele... java:117)  
at me.ele... IB\$56e34aaf6.invoke(<generated>)  
at org.springframework.cglib.proxy.MethodProxy.invoke(MethodProxy.java:201)

Thank You

Q&A